

In [2]:

```
1 import numpy as np
2 import pandas as pd
3
4
5 from sklearn.linear_model import ElasticNet
6 from sklearn.metrics import mean_squared_error
7 from sklearn.model_selection import train_test_split
8 from sklearn.model_selection import ShuffleSplit
9 from sklearn.model_selection import GridSearchCV
10 from sklearn.pipeline import Pipeline
11 from sklearn.preprocessing import StandardScaler
```

executed in 1.19s, finished 13:50:29 2020-05-22

**Формулировка на простом языке:**

Необходимо сделать предсказание количества совершаемых поездок в зависимости от погодных условий.

**Формулировка на математическом языке:**

Ставится задача регрессии для количества совершаемых поездок в зависимости от погодных условий.

In [3]:

```
1 DATA_DIR = "../data/processed/"
```

executed in 6ms, finished 13:50:29 2020-05-22

In [4]:

```
▶ 1 def get_data(DATA_DIR): ↵
```

executed in 149ms, finished 13:50:30 2020-05-22

Выберем RMSE в качестве метрики(имеет ту же размерность, что и исходные данные, в отличие от MSE).

In [5]:

```
▼ 1 def root_mean_squared_error(y_true, y_pred):
  2     return np.sqrt(mean_squared_error(y_true, y_pred))
```

executed in 95ms, finished 13:50:30 2020-05-22

In [6]:

```
1 data = get_data(DATA_DIR)
2 data.head(5)
```

executed in 51.0s, finished 13:51:21 2020-05-22

Out[6]:

	trips_counter	Max_Temperature_F	Mean_Temperature_F	Min_TemperatureF	Max_Dr
0	409	71	62.0	54	
1	491	63	59.0	55	
2	313	62	58.0	54	
3	395	71	61.0	52	
4	294	64	60.0	57	

In [7]:

```
1 data = pd.concat((data.iloc[:, :-1], pd.get_dummies(data.iloc[:, -1])),axis=1)
```

executed in 14ms, finished 13:51:21 2020-05-22

Попробуем применить линейную регрессию, но не простую, а Elastic Net.

In [8]:

```
1 X_train, X_test, y_train, y_test = train_test_split(data.iloc[:, 1:],
2                                                    data.iloc[:, 0],
3                                                    test_size=0.2,
4                                                    random_state=44)
```

executed in 80ms, finished 13:51:21 2020-05-22

Посмотрим на размер тренировочной/тестовой выборки.

In [9]:

```
1 X_train.shape, X_test.shape
```

executed in 5ms, finished 13:51:21 2020-05-22

Out[9]:

((551, 28), (138, 28))

In [10]:

```
1 model = ElasticNet(max_iter=5000)
2 steps = [('scaler', StandardScaler()), ('model', model)]
3 pipeline = Pipeline(steps)
```

executed in 5ms, finished 13:51:21 2020-05-22

In [11]:

```
1 parameters_grid_elastic = {  
2     'model__alpha' : np.linspace(0.00001, 2, num=100),  
3     'model__l1_ratio' : np.linspace(0, 1, num=50)  
4 }
```

executed in 6ms, finished 13:51:22 2020-05-22

In [12]:

```
1 ss = ShuffleSplit(n_splits=5, test_size=0.25, random_state=44)
```

executed in 3ms, finished 13:51:22 2020-05-22

In [13]:

```
1 gs = GridSearchCV(estimator=pipeline,  
2                   param_grid=parameters_grid_elastic,  
3                   scoring='neg_mean_squared_error',  
4                   cv=ss,  
5                   n_jobs=-1,  
6                   verbose=2)
```

executed in 5ms, finished 13:51:22 2020-05-22

In [14]:

```
1 gs.fit(X_train, y_train)
```

executed in 3m 1s, finished 13:54:23 2020-05-22

Fitting 5 folds for each of 5000 candidates, totalling 25000 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

```
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed:    3.6s
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed:    9.3s
[Parallel(n_jobs=-1)]: Done 872 tasks     | elapsed:   18.4s
[Parallel(n_jobs=-1)]: Done 3136 tasks    | elapsed:   33.4s
[Parallel(n_jobs=-1)]: Done 6056 tasks    | elapsed:   54.8s
[Parallel(n_jobs=-1)]: Done 9616 tasks    | elapsed:   1.3min
[Parallel(n_jobs=-1)]: Done 13832 tasks   | elapsed:   1.8min
[Parallel(n_jobs=-1)]: Done 18688 tasks   | elapsed:   2.3min
[Parallel(n_jobs=-1)]: Done 24200 tasks   | elapsed:   2.9min
[Parallel(n_jobs=-1)]: Done 25000 out of 25000 | elapsed:   3.0min finished
```

Out[14]:

```
GridSearchCV(cv=ShuffleSplit(n_splits=5, random_state=44, test_size=0.25, train_size=None),
             error_score=nan,
             estimator=Pipeline(memory=None,
                                 steps=[('scaler',
                                         StandardScaler(copy=True,
                                                         with_mean=True,
                                                         with_std=True
                                         e)),
                                         ('model',
                                          ElasticNet(alpha=1.0, copy_X=True,
                                                         fit_intercept=True,
                                                         l1_ratio=0.5, max_iter=5000,
                                                         normalize=False,
                                                         positive=False,
                                                         precompute=False,
                                                         random_state=0.51020408, 0.53061224, 0.55102041, 0.57142857, 0.59183673,
                                                         0.6122449 , 0.63265306, 0.65306122, 0.67346939, 0.69387755,
                                                         0.71428571, 0.73469388, 0.75510204, 0.7755102 , 0.79591837,
                                                         0.81632653, 0.83673469, 0.85714286, 0.87755102, 0.89795918,
                                                         0.91836735, 0.93877551, 0.95918367, 0.97959184, 1.
                                                         ])),
                                 pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                                 scoring='neg_mean_squared_error', verbose=2)
```

In [15]:

```
1 root_mean_squared_error(y_test, gs.best_estimator_.predict(X_test))
```

executed in 8ms, finished 13:54:23 2020-05-22

Out[15]:

75.65636238907558

In [1]:

```
▼ 1 def mean_absolute_percentage_error(y_true, y_pred):  
  2     return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
```

executed in 14ms, finished 13:50:20 2020-05-22

In [16]:

```
1 mean_absolute_percentage_error(y_test, gs.best_estimator_.predict(X_test))
```

executed in 11ms, finished 13:54:24 2020-05-22

Out[16]:

23.550110148977645

Получили результат близкий к тем, что мы получали у CatBoost/RandomForest, при этом MAPE оказывается меньше.